



Autoscaler Comparison: Custom HPA vs Kubernetes HPA


Incoming Request Rate (Workload)



p99 Server-Side Inference Latency



CPU Cores Allocated (= Replica Count)



Dropped Requests Rate (503s)

