

Start-of-Run vs End-of-Run Summary

Autoscaler	p99 Latency (start)	p99 Latency (end)	CPU Cores (start)	CPU Cores (end)
Custom HPA (queueing theory)	100 ms	100 ms	2	6
Kubernetes HPA 70 % CPU target	100 ms	100 ms	2	2
Kubernetes HPA 90 % CPU target	326 ms	100 ms	2	2